

1 Syntax of obfuscation

Definition 1.1 (Obfuscator Syntax). An obfuscator for a class of circuits $\mathcal{C}_{\lambda \in \mathbb{N}}$ is a randomized algorithm obf such that $\text{obf}(C, 1^\lambda)$ returns a functionally equivalent circuit to C for all $C \in \mathcal{C}_\lambda$, formally, for all $C \in \mathcal{C}_\lambda$, we have

$$\Pr_{\tilde{C} \leftarrow \text{obf}(C, 1^\lambda)} [\forall x \in \{0, 1\}^\ell : \tilde{C}(x) = C(x)],$$

where $\ell = \text{input} - \text{size}(C)$

2 Impossibility of Virtual Black-Box (VBB) Obfuscation

In 2001, Barak, Goldreich, Impagliazzo, Rudich, Sahai, Vadhan, and Yang (BGIRSY) showed that an obfuscator cannot provide security as good as having only black-box access to a circuit, known as *virtual black-box obfuscation*. Their main idea is to take a circuit describing a point function $f_{x,y}$ for x and y drawn uniformly from $\{0, 1\}^\lambda$ and a circuit describing a point testing function $t_{x,y,s}$ where additionally s is drawn uniformly random from $\{0, 1\}^\lambda$:

$\frac{f_{x,y}(x')}{\text{if } x = x'}$	$\frac{t_{x,y,s}(f)}{\text{if } f(x) = y}$
$\text{return } y$	$\text{return } s$
$\text{else return } 0^{ y }$	$\text{else return } 0^{ s }$

Now, given an obfuscation of $f_{x,y}$ and $t_{x,y,s}$, one can run $t_{x,y,s}$ on the obfuscation of $f_{x,y}$ and get s as a result. However, when only having black-box access to $f_{x,y}$ and $t_{x,y,s}$, this is not possible. Hence, an obfuscator must always leak more than black-box access. For a detailed discussion and the definition of virtual black-box obfuscation, see <https://eccc.weizmann.ac.il/eccc-reports/2001/TR01-057/index.html>.

3 Indistinguishability obfuscation

BGIRSY explored definitions of obfuscations which were not affected by their impossibility result. They found that *indistinguishability obfuscation* was not affected. 12 years later, Garg, Gentry, Halevi, Raykova, Sahai and Waters presented the first candidate indistinguishability obfuscation construction <https://eprint.iacr.org/2013/451>. Equally importantly, in the same month, Sahai and Waters presented their puncturable program technique which revealed the tremendous usefulness of indistinguishability obfuscation <https://eprint.iacr.org/2013/454>. Earlier this year, Jain, Lin and Sahai, in breakthrough result, based an indistinguishability obfuscator on widely believed assumptions, after 7 years of intense research in the field, see

<https://www.ias.edu/video/indistinguishability-obfuscation-well-founded-assumptions>

for a talk by Lin and <https://eprint.iacr.org/2020/1003> for the paper.

But which level of security does indistinguishability obfuscation actually provide? Indistinguishability obfuscation ensures that, given two circuits C_0 and C_1 *with the same functionality*, we cannot distinguish between obfuscations of C_0 and C_1 .

Definition 3.1 (Indistinguishability Obfuscation). Let p be a polynomial and obf be an obfuscator for $(\mathcal{C}_\lambda)_{\lambda \in \mathbb{N}}$, where $(\mathcal{C}_\lambda)_{\lambda \in \mathbb{N}}$ is such that for all $\lambda \in \mathbb{N}$, $|\mathcal{C}_\lambda| \leq p(\lambda)$. We consider a PPT algorithm \mathcal{S} an *equivalent circuit sampler* for $(\mathcal{C}_\lambda)_{\lambda \in \mathbb{N}}$ if for all $\lambda \in \mathbb{N}$, it holds that

$$\Pr_{(C_0, C_1) \leftarrow \mathcal{S}(1^\lambda)} \left[C_0, C_1 \in \mathcal{C}_\lambda \wedge \ell(C_0) = \ell(C_1) \wedge \forall x \in \{0, 1\}^{\ell(C_0)} : C_0(x) = C_1(x) \right],$$

where $\ell(C_0)$ denotes the input length of C_0 and $\ell(C_1)$ denotes the input length of C_1 . obf is *indistinguishable under equivalent circuit sampling* (IND-ECS) if for all equivalent circuit samplers \mathcal{S} for $(\mathcal{C}_\lambda)_{\lambda \in \mathbb{N}}$, it holds that $\text{Gind-ecs}_{\text{obf}, \mathcal{S}}^0$ and $\text{Gind-ecs}_{\text{obf}, \mathcal{S}}^1$ are indistinguishable, i.e., for all PPT adversaries \mathcal{A} , we have that

$$|\Pr[1 = \mathcal{A} \rightarrow \text{Gind-ecs}_{\text{obf}, \mathcal{S}}^0] - \Pr[1 = \mathcal{A} \rightarrow \text{Gind-ecs}_{\text{obf}, \mathcal{S}}^1]|$$

is negligible.

<u>$\text{Gind-ecs}_{\text{obf}, \mathcal{S}}^0$</u>	<u>$\text{Gind-ecs}_{\text{obf}, \mathcal{S}}^1$</u>
<u>Parameters</u>	<u>Parameters</u>
λ : security parameter	λ : security parameter
obf : obfuscator	obf : obfuscator
\mathcal{S} : equiv. circ. sampler	\mathcal{S} : equiv. circ. sampler
 <u>Package State</u>	 <u>Package State</u>
no state	no state
 <u>OBF()</u>	 <u>OBF()</u>
$(C_0, C_1) \leftarrow \mathcal{S}(1^\lambda)$	$(C_0, C_1) \leftarrow \mathcal{S}(1^\lambda)$
$C \leftarrow \text{obf}(C_0, 1^\lambda)$	$C \leftarrow \text{obf}(C_1, 1^\lambda)$
return (C_0, C_1, C)	return (C_0, C_1, C)

Remark. Note that the only difference between $\text{Gind-ecs}_{\text{obf}, \mathcal{S}}^0$ and $\text{Gind-ecs}_{\text{obf}, \mathcal{S}}^1$ is that one of the games obfuscates C_0 and the other obfuscates C_1 .